

Software Craftsmanship

Los Del DGIIM, losdeldgiim.github.io

Doble Grado en Ingeniería Informática y Matemáticas

Universidad de Granada

Esta obra está bajo una Licencia Creative Commons Atribución-NoComercial-SinDerivadas 4.0 Internacional (CC BY-NC-ND 4.0).

Eres libre de compartir y redistribuir el contenido de esta obra en cualquier medio o formato, siempre y cuando des el crédito adecuado a los autores originales y no persigas fines comerciales.

Software Craftsmanship

Los Del DGIIM, losdeldgiim.github.io

Arturo Olivares Martos

Granada, 2025-2026

Índice general

1. Clean Code	5
1.1. Meaningful names	5
1.2. Functions	6
1.3. Comments	7
2. VCS With Git	11
2.1. Git Data Model	11
2.2. Working with Others	12
3. Test Driven Development	13
3.1. Debug-Later Development	13
3.2. Test Driven Development	14
3.2.1. The Three Laws of TDD	14
3.2.2. The TDD Microcycle	14
3.3. Testing	15
3.3.1. Red Bar Patterns	15
3.3.2. Green Bar Patterns	16
3.3.3. Testing Patterns	16
4. Agile Development	19
4.1. Agile Methodologies	20
4.1.1. Scrum	20
4.1.2. XP (Extreme Programming)	21
5. Mid-level Software Design	25
5.1. The SOLID Principles	26
5.1.1. Single-Responsibility Principle (SRP)	26
5.1.2. Open-Closed Principle (OCP)	27
5.1.3. Liskov Substitution Principle (LSP)	28

1. Clean Code

We will focus on code, as we don't only want to write good code, but also tell the difference between good and bad code. However, we should firstly assume some claims:

- There will always be code

In a system, there will always be requirements (as with no requirements, the system would not be needed). Code is simply a high specification of those requirements. In addition, given that we will have programming languages with high abstraction levels and that we will use AI, code will not disappear.

- Code is written for humans

Code is almost never just written once and then forgotten. We should always assume that code will be read by other people, and thus, we should write it in a way that is easy to understand.

- Bad code will bring your company down

If bad code is written, the cost of solving bugs will increase, and thus, the company will lose money. Code smells, which are indicators of bad code (they may not always be bugs), may include duplicated or commented code, long methods, or large classes.

- Start from the scratch will not fix the problem

When bad code becomes so unmanageable, it may be tempting to start from scratch. However, this is not a good idea, as it will take a lot of time and resources, it may not solve the problem, and it may even lead to two systems developed in parallel, which will be a nightmare to maintain.

In order to solve the problem, code should always be kept clean, following the *Boyscout Rule*: "Always leave the code cleaner than you found it".

Once we have assumed these claims, we can start talking about clean code. Even though there is no single definition of clean code, it should just be as easy, minimal and *simple* as possible. It should clearly express its intent. We will focus on some points that will help us to write clean code.

1.1. Meaningful names

We name a lot of components in our code, such as variables, functions, classes, etc. We should always try to give them meaningful names, as they will help us to understand the code.

1. Use intention-revealing names.

`int x;` is not a good name, as it does not reveal the intent of the variable.

2. Avoid disinformation.

We should not misuse known abbreviations, imply data types, or use similar names for different things.

3. Make meaningful distinctions.

We should not use the same name for different things, as it will be confusing. What is the difference between `a1` and `a2`? We should use names that clearly distinguish between different things.

4. Avoid encodings

Even though it was done in the past, type of scope information should not be encoded in the name, as it should nowadays be provided by the IDE.

5. Regarding classes:

Nouns should be used for classes, while verbs should be used for methods. Accessors (**getters**), mutators (**setters**), and predicates should be named accordingly, as they are very common and should be easily identifiable.

6. Solution vs Problem Domain names

Code that is related to the problem domain should be named accordingly, while code that is related to the solution domain should be named in a way that reflects its purpose.

7. Use short, pronounceable names.

A name should be as short as possible, as long as it provides enough context.

Finding good names is not easy, but it is worth the effort, as it will make our code much easier to understand and maintain.

1.2. Functions

Functions are one of the most important components of our code, as they are the building blocks of our programs. We should always try to write functions that are clean and easy to understand. Some guidelines for writing clean functions are:

1. Small functions.

Functions should be small (ideally, 3 lines or less), as they are easier to understand and maintain. If a function is too long, it may be a sign that it is doing too much and should be split into smaller functions.

This may lead to a lot of functions, aspect that may be criticized and that may even impact performance. Therefore, each one should find the right balance between the number of functions and their size.

2. Do one thing.

A function should do one thing and do it well (Error handling is needed and is not considered as doing more than one thing). If the developer is tempted to write a comment before a block of code, it may be a sign that the function is doing too much and should be split into smaller functions.

3. One Level of Abstraction per Function.

A function should not mix different levels of abstraction. Each time a function is extracted, it should be at a lower level of abstraction than the one that calls it.

4. Regarding arguments:

a) Low number of arguments.

Functions should have a low number of arguments (ideally, 0 or 1), as they could be messed up when calling the function. If too many arguments are needed, maybe they should be encapsulated in an object.

b) Output arguments should be avoided.

It is assumed that arguments flow in the function, not out of it.

c) Flag arguments should be avoided.

5. They should have no side effects.

6. Command / Query Separation.

Functions should either do something (command) or answer something (query), but not both. If a function does both, it may be a sign that it is doing too much and should be split into smaller functions.

7. DRY (Don't Repeat Yourself).

8. Exceptions should be used instead of return codes.

9. Multiple `return` statements are acceptable.

Functions should ideally follow a Strict-Structured Programming style, which means that they should have a single entry and a single exit point (only one `return` statement, no `break`, `continue`, or `goto` statements). However, in some cases, it may be acceptable to have multiple return statements, as long as they are used in a way that does not make the code harder to understand.

1.3. Comments

Comments are a necessary evil. They are really extendedly used, but they are not always good. They may be outdated, they may be misleading, and they may even be used to explain bad code. Bad comments lead to people ignoring them, and thus, they may even hide important information.

The programmer should explain himself through the code, not through comments. Some bad examples are:

- `int d; // elapsed time in days`
- Noise/Redundant Comments, as
`int calculatePay() { // calculate the pay`
- Using a difficult condition and then explaining it with a comment, instead of just using a bool variable with a meaningful name or a function that encapsulated the condition.
- Comments that try to explain bad code, instead of just refactoring it to make it easier to understand.
- Mandatory comments, which are required by the company or by the project, but that do not provide any useful information.
- Non-Local information

Comments should not provide information that is not local to the code they are commenting, as it may be easily forgotten and thus, lead to confusion.

- Comments of different sections of a function

If a function is too long and it is divided into different sections, it may be a sign that the function is doing too much and should be split into smaller functions.

- Closing brace comments

When a function is too long, it may be difficult to know which closing brace corresponds to which opening brace. However, if the function is too long, it should be split into smaller functions, and thus, closing brace comments should not be needed.

- Journal comments & Attributions

The VCS should be used to track changes and authorship, not comments.

- Commented-out code

Again, the VCS should be used to track changes and to recover old code if needed, not comments.

However, there are some good comments, such as:

- Legal comments, which are required by law or by the company.

- Explanation of intention

Sometimes, a decision is made in a way that is not obvious, and it may be worth explaining the intention behind it. People may not agree with the decision, but at least they will understand it.

- Warning of consequences

- Amplification of a point in the code that is not obvious.

People may not understand why a certain point in the code is important, and it may be worth amplifying it with a comment. If not done, another developer may change that point in the code, not understanding its importance, and thus, breaking the code.

- Comments to create Public API documentation (e.g., Javadoc).

Lastly, there are some comments that are not good, but that may be necessary, such as:

- Clarification of a point in the code that is not obvious.

Sometimes, a point in the code is not obvious, and it may be worth clarifying it with a comment. For instance, when a function returns a value that is not obvious, or when a complex regular expression is used, it may be worth clarifying it with a comment. However, it is risky because if the code is changed, the comment may become outdated and thus, misleading.

- TODO comments, which indicate that something needs to be done in the future.

They may be used to ask a colleague to work on something, a reminder to change a part that depends on something else... but never to clean up bad code later.

Modern IDEs have tools to track TODO comments, but they should be used with caution, as they may be forgotten and thus, lead to technical debt. The Boy Scout Rule should be followed.

On conclusion, comments should be used with caution, as they may be outdated, misleading, and even hide important information. The code should explain itself, and if a comment is needed to explain a point in the code, it may be a sign that the code is not clean and should be refactored to make it easier to understand.

2. VCS With Git

In this chapter, we will focus on the insights of using Git as a Version Control System (VCS). However, a deeper understanding of the concepts behind Git can be found in the Chapter 2 of the contents of the [Application Management](#) course.

Definición 2.1 (Version Control System). A Version Control System (VCS) is a software tool that helps manage changes to source code over time. It allows multiple developers to collaborate on a project, track changes, and maintain a history of modifications.

Git is a distributed version control system that provides a powerful and flexible way to manage code changes. There is a lot of reasons why Git is widely used in the software development industry, including its speed, efficiency, and support for branching and merging. Even when someone works alone on a project, using Git can provide benefits such as keeping a history of changes, allowing to revert to previous versions, and facilitating the integration of new features or bug fixes.

2.1. Git Data Model

Git is based on three different types of objects: blobs, trees, and commits. All of these objects are identified by a unique SHA-1 hash, which is generated based on the content of the object, and saved in the `.git/objects` folder.

1. Blobs: A blob (binary large object) is a file that contains the contents of a file in the repository. It does not contain any metadata, such as the file name or permissions.
2. Trees: A tree is a directory that contains references to blobs and other trees. It represents the structure of the repository at a given point in time. It also contains metadata, such as the file name and permissions.
3. Commits: A commit is a snapshot of the repository at a specific point in time. It contains a reference to a tree, which represents the state of the repository at the time of the commit. It also contains metadata, such as the author, date, and commit message.

In order to refer commits in a more human-readable way, Git uses references, such as branches and tags. The most important ones are:

- **HEAD**: A reference to the current commit that the user is working on.

- **master or main:** The default branch in Git, which typically represents the main development line.

Commits always have a comment. There are good and bad practices for writing commit messages. A good commit message should be concise and should describe the changes made in the commit.

Another important concept in Git is commit granularity. Although some people have a different opinion on this topic, it is generally recommended to have small and focused commits that represent a single logical change. This makes it easier to understand the history of the repository and to revert changes if necessary.

In order to save the history of changes, Git uses a DAG (Directed Acyclic Graph) data structure. Each commit points to its parent commit(s), creating a chain of commits that represents the history of the repository. This allows Git to efficiently manage and track changes over time, enabling features like branching and merging.

Lastly, the idea of branches in Git is fundamental to its workflow. A branch is a pointer to a specific commit, allowing developers to work on different features or bug fixes without affecting the main codebase. Branches can be easily created, merged, and deleted, making it a powerful tool for managing parallel development efforts.

2.2. Working with Others

Git is a distributed VCS, so it uses a centralized collaboration. This means that there is a central repository that serves as the main source of truth for the project, and developers can clone this repository to their local machines, make changes, and then push those changes back to the central repository. This central repository can be hosted on platforms like GitHub, GitLab, or Bitbucket, which provide additional features for collaboration, such as pull requests, code reviews, and issue tracking.

3. Test Driven Development

This chapter covers the concept of *Testing*. Given that the goal of a developer is achieving clean code that *works*, testing is a crucial part of the development process. There are however two main approaches to testing: *Debug-Later Development* and *Test Driven Development*. The first one is the most common, but the second one is the one that leads to better code quality. We will see in this chapter why that is the case, and how to implement TDD in practice.

3.1. Debug-Later Development

As the name suggests, Debug-Later Development is the approach where developers write code without testing it, and then debug it later when it doesn't work. The problem is that a new bug may be found at any time, forcing us to go back and forth between writing code and debugging it. In order to explain this concept, the following time definitions are used:

- Time till discovery (T_d): the time it takes for a bug to be discovered since it was introduced.
- Time till found (T_{find}): the time it takes for the reason for the bug to be found since it was discovered.
- Time till fixed (T_{fix}): the time it takes for the bug to be fixed since its reason was found.

This approach has several disadvantages:

- T_d and T_{find} are usually very long, because the code is not tested until it is complete, and the bug may be found at any time.
- The longer a bug is present in the code, the more likely it is to cause other bugs and the harder it is to fix it. Therefore, T_{fix} is also usually very long.
- Clean code is harder to achieve, because developers are more focused on making the code work than on making it clean.
- Testing is usually forgotten or neglected, as developers think that their code works, and they don't want to spend time writing tests for it.

This is the reason why Debug-Later Development is not a good approach, and why Test Driven Development is a better alternative.

3.2. Test Driven Development

Test Driven Development (TDD) is a development discipline (not a testing technique) where developers write tests before writing the code that makes the tests pass. It helps sequentially focus individually on working code and clean code, and it leads to better code quality.

3.2.1. The Three Laws of TDD

The Test Driven Development process is based on three laws, which guarantee that the code is always tested:

- First Law: You are not allowed to write any production code unless it is to make a failing unit test pass.
- Second Law: You may not write more of a unit test than is sufficient to fail, and not compiling is failing.
This way, as soon as the new test fails, you must stop writing the test and start writing the production code to make it pass. In addition, as there is no production code yet, not compiling will happen very soon.
- Third Law: You may not write more production code than is sufficient to pass the currently failing test.

This leads to three main benefits. Firstly, only what is strictly necessary is coded. Secondly, bugs are found and fixed as soon as they are introduced (T_d , T_{find} and T_{fix} are reduced), which makes them easier to fix. Finally, tiny steps are always made towards the goal, which makes it easier to achieve clean code.

3.2.2. The TDD Microcycle

The TDD process can be summarized in a microcycle, which is a sequence of steps that are repeated every few minutes until the code is complete:

1. Write a small test (as small as possible).
2. Compile all the tests and see that the new test does not compile.
3. Write the production code to make the test compile.
4. Compile all the tests and see that the new test fails.

We should not assume that the test fails, as it may pass for the wrong reason (for example, if the test is not written correctly). If we assume that the test fails, we may code with a false sense of security.

5. Write the production code to make the test pass.
6. Compile all the tests and see that they *all* pass. It is important to check that all the tests pass, because the new code may have broken some existing code.
7. Refactor the code in order to make it cleaner.

8. Compile all the tests again and see that they all pass.

Regarding refactoring, it is the process of changing the code without changing its behavior. It can be risky, as it can introduce new bugs, so it should not be done until all the tests pass.

3.3. Testing

As explained before, testing is a crucial part of the development process, as they help us feel confident that our code works and it reduces the developers' stress. However, they are not easy at all, as they require a lot of practice and experience to be good at writing tests. In addition, it should be noted that tests *are also code*. They must be designed, written, refactored and maintained just like production code.

3.3.1. Red Bar Patterns

A red bar represents a failing test; so a red bar pattern is a pattern that leads to a failing test; which is what should be done according to the microcycle of TDD. There are some common red bar patterns that should be followed:

- List of Future Tests: a list of tests that should be written in the future should be written. It lets us focus on one test at a time, and it helps us keep track of the tests that we want to write. It should not be a fixed list, as it may change as we write the code and discover new things.
- Picking the Next Test: One that can easily be made to pass should be picked. This way, we can make progress quickly and feel good about it, which motivates us to keep going. If none applies, an existing bigger test can be split into smaller tests (which should also be added to the list of future tests).
It is also important to note that a passing test provides an assertion. Therefore, tests that allow to assume things that will be useful for future tests should be picked, as they will make it easier to write the future tests.
- Adding Tests: Test should be added to the list of future tests for several reasons: when a potential problem is introduced, when a bug is reported, when a big test is split into smaller tests... However, it should not be forgotten that the list may eventually be finished.
- Removing Tests: The list of future tests should be reviewed regularly, and tests that are no longer relevant should be removed, as more tests does not necessarily mean better testing.

If a test can be made fail individually, it should be kept, as it tests something that is not tested by other tests. However, if a test cannot be made fail individually, it may be testing something that is already tested by other tests, and its deletion may be considered.

There are two more aspects that should be taken into account regarding red bar patterns: what to do when stuck, and how to end the programming session.

- When Stuck: if we are stuck, our first attempt should be to take a short break, as it may help us clear our mind and see the problem from a different perspective. If we are still lost after the break, throwing away the code and starting over may be a good idea, as it allows us to start fresh and avoid getting stuck in a dead end.

Another good idea is to switch partners in pair programming, as it forces us to explain our code to the new partner, which may help us see the problem from a different perspective. In addition, the new partner is not emotionally invested in prior decisions, which may make it easier to throw away the code and start over if necessary.

- Ending the Programming Session: when it comes to ending the programming session:

- If we program alone, we should end the session with the last test failing, so that we know where to start the next session.
- If we program in teams, we should end the session with all the tests passing, so that we can leave the code in a good state for the next person.

3.3.2. Green Bar Patterns

A green bar represents a passing test; so a green bar pattern is a pattern that leads to a passing test. There are some common green bar patterns that should be followed:

- Obvious Implementation: the simplest implementation that makes the test pass should be written.
- Fake it ('til you make it): if the test is hard to make pass, a fake implementation (maybe returning a hardcoded value) can be written to make the test pass, and then it can be refactored to a real implementation.

However, it is important to make sure that you have a test that will force you to look at the fake implementation again, if not, you should add it to your list of future tests.

- Triangulation: if you don't know how to code the implementation, wait until you have two or more tests. In that case, you should use the *Rule of Thumb*: fake it until it is more trouble to fake it than to make it. Therefore, when you can no longer fake it, you should code the implementation, as you will have enough information to do it.

3.3.3. Testing Patterns

Some good testing principles that should be followed are:

- Small Tests: tests should be small. A test should fail for one reason, and it should be easy to understand why it fails.

- Always rerun all the tests: all the tests should be rerun after every change, as we may have broken something without realizing it.
- Tests should fail: if a bug is intentionally introduced, a test should fail.
- Tests should fail independently: ideally, each test should be able to fail independently of the others. If a lot of tests fail at the same time, it may be hard to understand why they fail and to fix the problem. It could also mean that the tests are not well designed, as they may be testing too much at the same time, or they may be redundant.
- Do not debug: in case of a failing test, if we can fix it without debugging, we should do it, as it is faster and it helps us keep the flow. However, if we cannot fix it directly, it is faster to go back (using VCS).
- Tests for external software: if we are using external software (for example, a database), we should write tests for it, as it may have bugs that affect our code.
- Test Private Methods through Public Methods: private methods should not be tested directly, as they are implementation details that may change. Instead, they should be tested through public methods that use them.
- Test Error Cases: tests should not only cover the correct cases, but also the errors. It should be checked that the errors are reported and handled correctly. With this aim, crash test dummies (explained in the following subsection) can be used.
- Test Should be Fast: tests should be fast, as they will be run frequently. If tests are slow, we may be tempted to run them less frequently (which is not desired). Possible ways to make tests faster are mocking external software (explained in the following subsection) or removing unnecessary tests.
- Test Unit Behavior: tests should not necessarily test the whole class, but they should test the behaviors of the unit. The tests' name should reflect the behavior that they are testing.
- Legacy Code: if we have legacy code (code without tests), there are two possible approaches to add tests to it:
 - If you can add unit tests directly, do it but just testing its current behavior (not caring about if it's correct or not).
 - If you cannot add unit tests directly, start from the outside to the inside: start from the outside (whole system tests), use them to write integration tests, and then use them to write unit tests.

This way, you will have a safety net to refactor the code and add more tests later.

There is one more aspect that should be taken into account, mocking.

Mocking

Mocking is the process of replacing a real object with a fake one that simulates its behavior. It is useful when we want to test a unit that depends on an external software (for example, a database), or when we want to test a unit that is not yet implemented.

Mocking is in fact a *slang* term for *test double*, which is a more general term that includes different types of fake objects:

- Test Dummy: an object that is passed around but never actually used (only used to satisfy the compiler).
- Test Stub: an object that provides predefined answers to method calls.
- Crash Test Dummy: special stub that return errors (throws exceptions) when called. It is mainly used to test error handling.
- Exploding Fake: an object that causes test failure if it is called. It is mainly used to test that a method is not called when it should not be.
- Test Spy: special stub that records information about how it was called, which functions were called, with what parameters, how many times...
- Mock Object: more complex spy that also includes assertions about how it should be called.
- Fake Object: an object that has a working implementation, but it is not suitable for production (for example, an in-memory database).

When injecting a test double, they should not be injected directly, as it may make the code more complex and less maintainable. Instead, they should be injected as a reference to the class. Therefore, the test double can be passed when testing, and the real object can be passed when running the code in production.

Mocking has several advantages, as it allows us to test units that depend on external software, and it allows us to test units that are not yet implemented. It also allows for faster tests. Therefore, a common mistake is thinking that every dependency should be mocked, but that can lead to slower tests and explosion of polymorphism (for dependency injection). Therefore, the “Uncle Bob’s Mocking Heuristic” should be followed: mock across architecturally significant boundaries, but not within those boundaries.

4. Agile Development

Software Development has always been a field in constant evolution, with new methodologies and practices emerging to address the challenges of creating high-quality software efficiently. One of the most influential approaches in recent years has been Agile Development, which emphasizes flexibility, collaboration, and customer satisfaction.

There are three main phases in the history of Software Development:

1. Pre-Agile Era:

During the early days of software development, it was done in small controllable steps, with a focus on craftsmanship and quality. Projects were smaller, and developers had more experience.

2. Industrialization: Scientific Management:

From the 1970s, the most common approach was the Waterfall model, which is a linear and sequential process. Coding should be postponed until the design phase is complete, and everyone should follow the plan to the letter. This is ideal for projects with high cost of change and well defined problems, which is actually not really common.

3. Reformation: the comeback of Agile:

It started in the 1990s with the emergence of new methodologies like Scrum and Extreme Programming (XP). In 2001, a group of software developers created the Agile Manifesto, which outlined the core values and principles of Agile Development.

In typical development process, after the common phases (requirements gathering, design, implementation...) usually a well-known phase arises: the *Death march phase*, where the project is in a critical state, with a lot of pressure to deliver, and the team is working overtime to meet the deadlines. This happens often because they have no real data that informs about the actual progress of the project, but just the *hope* to meet the deadlines. Agile provides data before the hope can kill the project.

In Figure 4.1 we can see the Iron Cross, which represents four qualities that software projects should have: *fast, cheap, good, done*. However, it is said that you can only have three of these qualities at the same time, and you have to sacrifice one of them. In order to correctly manage the iron cross, data is needed. The main core of Agile is to provide these data.

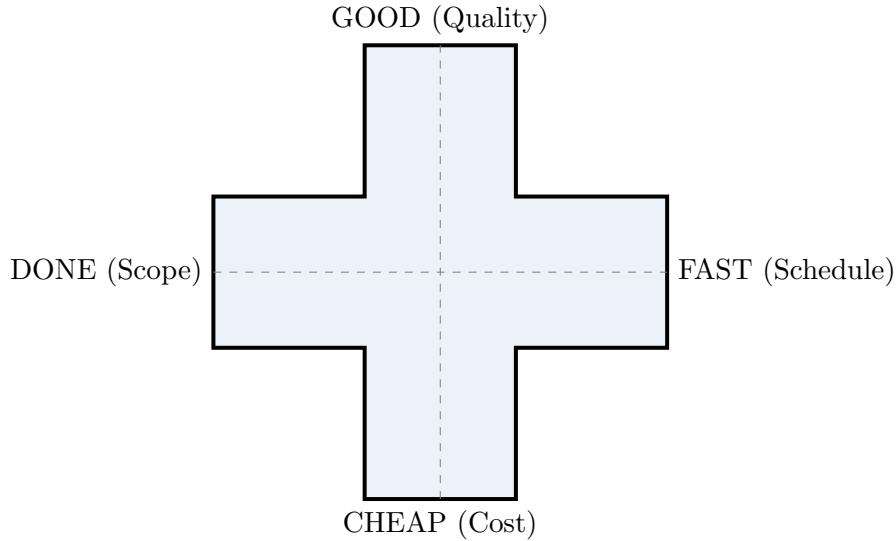


Figura 4.1: The Iron Cross of Software Development

There are two main questions whose answer (based on data, not on hope) is the main goal for Agile:

- *Do we do the right thing?*: Fast iterations, early feedback and planning are used.
- *Can we still do it in time?*

Apart from that, some core Agile values are:

- Personal courage and risk taking.
- Intense communication and collaboration, ignoring barriers of hierarchy, location and time.
- Rapid feedback and learning.
- Simplicity of code, design and teams.
- Respect.

There are different Agile methodologies, such as Scrum, Extreme Programming (XP), Crystal... However, they all share the same core values and principles, and they all aim to provide data to manage the iron cross of software development.

4.1. Agile Methodologies

4.1.1. Scrum

Scrum is one of the most popular Agile methodologies, and it is based on the idea of iterative and incremental development. A deeper understanding of the concepts behind Scrum can be found in the Chapter 1 of the contents of the [Application Management](#) course.

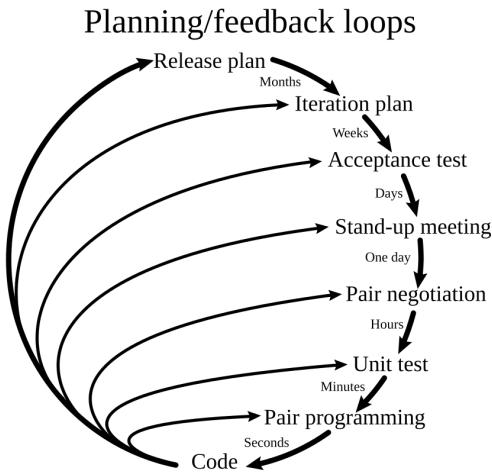


Figura 4.2: Planning and feedback loops in XP

4.1.2. XP (Extreme Programming)

XP is another popular Agile methodology, and it is based on the idea of continuous feedback and improvement. In the Figure 4.2 we can see that it uses a planning and feedback loop that represents how frequent are the feedbacks and iterations in XP.

It is based on a concept named “Circle of Life”, which represents the practices that should be taken, from a bigger aspect (the whole project) to a smaller aspect (the code itself).

Business-facing practices

They are practices whose aim is to reduce the common gap between the business and the development team. The main goal here is planning correctly, and therefore a project should be divided into small pieces, where each piece should be estimated individually. This estimation should be as accurate as possible but only as precise as needed. Some techniques for estimation are:

- Trivariate Analysis: Long-term estimation technique.

How much time will it take to do a task should be estimated three times:

- *Worst case*: Time within which the task will be completed with a 95 % confidence.
- *Nominal case*: Time within which the task will be completed with a 50 % confidence.
- *Best case*: Time within which the task will be completed with a 5 % confidence.

- Story Points: Short-term estimation technique.

The project is divided in stories.

- An user story is an abbreviated description of a feature from the perspective of the user. It typically follows the format: “As a [type of user], in order to [achieve a goal], I will [perform an action]”.

The stories should be made between the developpers and the stakeholders, and the details should be included when started working on the story, not before.

Each story should be estimated in story points.

- A story point is a unit of measure for expressing the overall effort (not time) required to implement a user story.

An initial story (called the “golden story”) of average complexity is chosen as a reference, and the other stories are estimated relative to it. For example, if the golden story is estimated to be 5 story points, and another story is estimated to be twice as complex, it would be assigned 10 story points. This technique allows for more accurate estimation and better planning of the project.

Then, the project evolves in iterations (each one finishing with a demo), which are organized in interation planning meetings.

- In iteration planning meetings, the stakeholders select a set of user stories to be implemented in the next iteration, based on their priority and the team’s *velocity* (the number of story points the team can complete in an iteration).

The initial velocity is an initial guess, and the velocity of next iterations is calculated based on the velocity of the previous iteration. Managers may have a velocity chart, which is a visual representation of the team’s velocity over time, to help track progress and make informed decisions about future iterations.

In addition, a *midpoint check* is usually done at the middle of the project, where the team reviews the story points completed so far, and compares it with the initial estimation. This allows to adjust the remaining story points asking the stakeholders to add/remove stories.

Team-facing practices

They are practices whose aim is to improve the communication and collaboration within the team. Some of these practices are:

- Metaphore / Ubiquitous language: A common language that is used by all members of the team, including developers, testers, and stakeholders.
- Sustainable pace
- Collective ownership
- CI
- Standup meetings

Technical practices

They are practices whose aim is to improve the quality of the code and the design of the software. Some of these practices are:

- TDD
- Pair programming: Two developers work together on the same code, sharing the workstation (it can also be done virtually). One developer writes the code, while the other reviews it in real-time. This practice promotes collaboration, knowledge sharing, helps to catch errors early in the development process, and also helps to train new team members. Pairing of larger teams can also be done, which is called mob programming.

Some characteristics that pair programming should have are:

- Optional
- Intermittent
- Unscheduled
- Short-lived
- Refactoring
- Simple design

Only the simplest design that works should be implemented. Some rules for simple design, in this priority order, are:

1. Pass all tests.
2. Reveal the intention of the code.
3. Remove duplication.
4. Decrease the number of elements (classes, methods, variables...).

5. Mid-level Software Design

In this course, our aim is to design good quality software, and to do that we need to know how to design it. The following are smells of bad design:

- Rigidity: the software is hard to change because every change affects too many other. Everything that *depends* on the changed thing needs to be changed too.
- Fragility: when you change something, it breaks something else that you didn't expect to break. There are hidden *dependencies*.
- Immobility: the software is hard to reuse because it's too entangled with the current application. The desired design is highly *dependent* on undesired parts of the software.

As the reader may have noticed, the common main root cause of all these smells is *dependencies*. Therefore, the solution to these problems is to design software with low dependencies. In order to do so, keeping a *simple* design is key. Only the code that is required should be written, with a structure that keeps it simplest, smallest and most expressive. In order to do so, the following rules should be followed (in this order):

1. Pass all the tests
2. Reveal the intent
3. Remove duplication
4. Decrease elements of the software

In this chapter, we will focus on the mid-level design of software. There are 4 levels:

1. Functions: already covered in Section 1.2.
2. Classes
3. Components
4. System or Architecture

Therefore, in this chapter we will cover the design of classes.

Definición 5.1 (Class). Coupled grouping of functions and data.

Observación. Even though that is the definition, there are a lot of programming languages that call *classes* to things that are not classes; and also there are a lot of programming languages that give other names to things that are classes (for instance, *structs* in C++).

Regarding the dependencies mentioned before, and given that classes are connected to others, there are two types of dependencies:

- Compile-time dependencies: they are the ones that are created when the software is being compiled. They are produced when one class mentions another class (or its functions or data) in its code.
- Runtime dependencies: they are the ones that are created when the software is running. They are produced when one class transfers control to another's functions or reads/writes data from another class.

One example of runtime dependency that is not a compile-time dependency is the case of *interfaces* (or *abstract classes*). A class *A* may have a compile-time dependency on an interface *I*, but then at runtime it may have a runtime dependency on a class *B* that implements *I*, without having a compile-time dependency on *B*.

During the following section, we will cover some important principles for designing good classes.

5.1. The SOLID Principles

The SOLID principles are a set of five design principles that are intended to make software more maintainable and flexible. Applying them too often/early can lead to over-engineering and unnecessary complexity, but they are very useful when applied in the right way.

5.1.1. Single-Responsibility Principle (SRP)

A class should have only one reason to change.

If a class has more than one responsibility (reason to change), they become coupled, so changes in one responsibility may affect the other. Ideally, each responsibility should be encapsulated in a separate class, but if not, at least they should be separated in different interfaces.

The SRP may lead to *temporary duplication*, as the same code may be needed in two different classes, but if it is due to different responsibilities (reasons to change), then this duplication may eventually disappear as the two classes evolve in different ways.

Another common symptom of a violation of the SRP is the presence of *merge conflicts*, as two different developers may be working on the same class for different reasons, and therefore merge conflicts may appear.

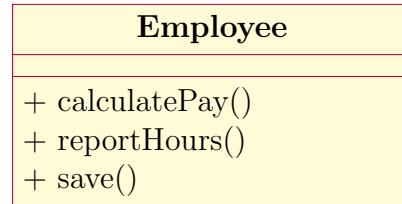


Figura 5.1: Example of a class that violates the Single-Responsibility Principle.

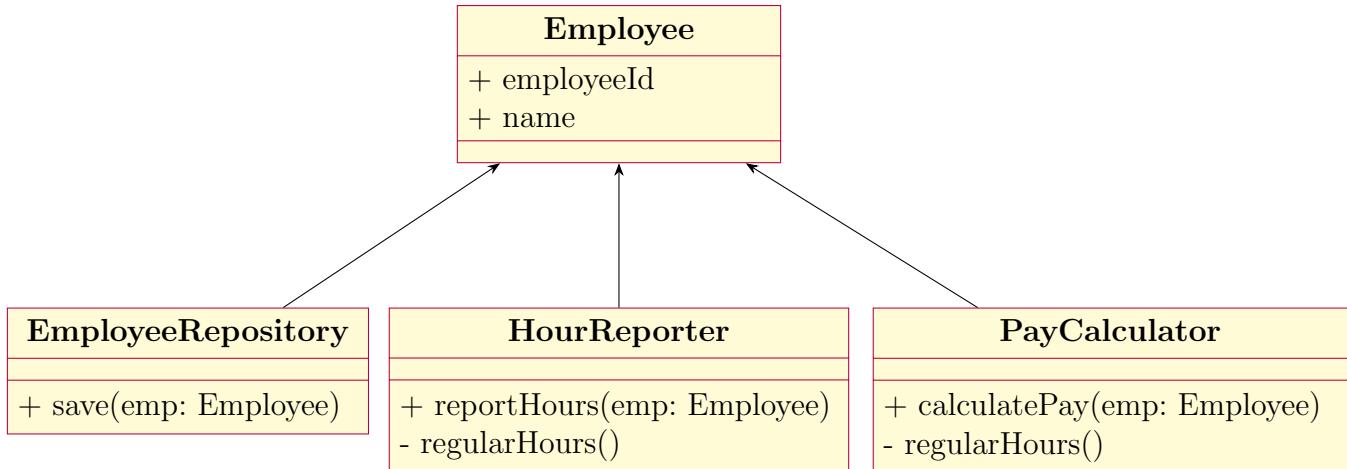


Figura 5.2: Example of a design that follows the Single-Responsibility Principle.

One example of a violation of the SRP is the class **Employee** described in Figure 5.1. This class has three responsibilities, which may lead to three different reasons to change. A possible solution to this problem is to split the class into three different classes, each one with a single responsibility, as shown in Figure 5.2. The three responsibilities are:

- Calculating the pay of the employee (possibly used by the finance department).
- Reporting the hours worked by the employee (possibly used by the HR department).
- Saving the employee data (possibly used by the IT department).

5.1.2. Open-Closed Principle (OCP)

Software entities (classes, modules, functions, etc.) should be open for extension, but closed for modification.

They should open for extension, as new requirements may mean new behavior that needs to be added to the software. However, they should be closed for modification, as changing existing code may introduce bugs and break existing functionality.

An example of a violation of the OCP is the design shown in Figure 5.3, as trying to connect the **Client** to a new **NewServer** would require modifying the **Client** class. A possible solution to this problem is to introduce an interface **Connection**



Figura 5.3: Example of a design that violates the Open-Closed Principle.

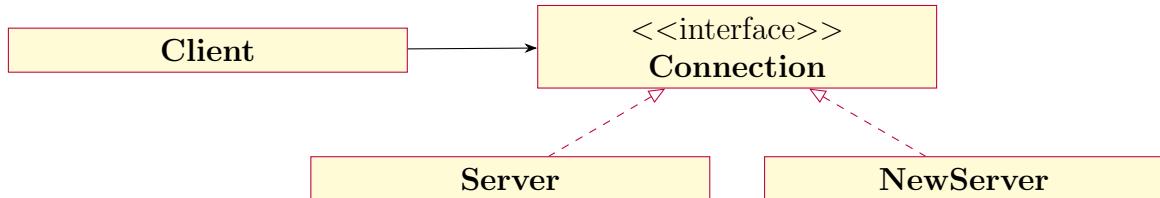


Figura 5.4: Example of a design that follows the Open-Closed Principle.

that both **Server** and **NewServer** implement, as shown in Figure 5.4. This way, the **Client** can connect to any class that implements the **Connection** interface without needing to be modified.

There will always some kind of change against which a module is not closed, so deciding which changes to be closed against is a matter of judgment. As many data as available should be collected (asking customers or domain experts, using personal experience, etc.) to make an educated guess. A common practice is considering that no changes will happen and, when a change happens:

1. We first implement the abstractions needed to make the change possible.
2. Then we implement the change itself.

5.1.3. Liskov Substitution Principle (LSP)

Subtypes must be substitutable for their base types.